Mini Project Report

on

CLARITY - Unmasking Political Question Evasions

Submitted by

Arpenaboyina Rakesh (22BCS016)

Azmeera Sai (22BCS023)

Chidari Sai Krishna (22BCS030)

Tanuku Venkata Vishnu Sai Karthik (22BCS129)

Under the guidance of

Dr. Sunil Saumya

Assistant Professor

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

INDIAN INSTITUTE OF INFORMATION TECHNOLOGY, DHARWAD

14/11/2025

Certificate

This is to certify that the project, entitled **CLARITY - Unmasking Political Question Evasions**, is a bonafide record of the Mini Project coursework presented by the students whose names are given below during 2025–2026 in partial fulfilment of the requirements of the degree of Bachelor of Technology in Computer Science and Engineering.

Roll No	**Names of Students**
22BCS016	Arpenaboyina Rakesh
22BCS023	Azmeera Sai
22BCS030	Chidari Sai Krishna
22BCS129	Tanuku Venkata Vishnu Sai Karthik

Dr. Sunil Saumya
(Project Supervisor)

Contents

List of Figures

List of Tables

1 Introduction

In recent years, the analysis of clarity and ambiguity in political communication has become an increasingly significant area of research within computational linguistics and political discourse analysis. Political interviews, press conferences, and debates are often characterized by evasive, indirect, or ambiguous responses, a phenomenon broadly known as equivocation. Equivocation occurs when speakers respond to questions in a way that appears relevant but lacks directness or explicitness—an intentional or strategic form of ambiguity.

Understanding and automatically detecting equivocation or unclear responses has both linguistic and societal importance. From a linguistic standpoint, it offers insight into how meaning and intent are communicated indirectly; from a political standpoint, it serves as a tool for accountability and transparency, allowing analysts and the public to assess whether a speaker genuinely answers for a particular question.

The study of equivocation in political communication dates back several decades. Classical frameworks by Bull (2003), Harris (1991), and Rasiah (2010) categorized responses as *Replies, Non-Replies,* and *Intermediate Replies* depending on how completely the speaker addressed the question. These works also introduced subtypes of evasive behavior, such as *deflection*, *dodging*, *attacking the question*, and *declining to answer*.

However, earlier typologies were primarily qualitative, focusing on manual annotation of small corpora and lacking large-scale computational validation. As a result, such analyses could not be generalized across diverse contexts or tested using modern machine learning methods.

Recent advances in Natural Language Processing (NLP) and Large Language Models (LLMs) have enabled researchers to revisit this problem computationally. Transformer-based architectures like **BERT**, **RoBERTa**, and **DeBERTa** can represent the semantic structure and contextual dependencies of questions and responses, allowing nuanced judgments about response clarity, ambiguity, and evasion.

The dataset contains 3,445 question–answer pairs from political interviews of U.S. presidents (2006–2023), annotated at two levels:

- **High-Level Clarity Classification:** determining whether the response is a
  - **Clear Reply** — explicit and direct
  - **Ambivalent Reply** — indirect or multi-interpretive

  - **Clear Non-Reply** — explicit refusal or denial

- **Low-Level Evasion Classification:** identifying one of nine fine-grained evasion techniques, including *Deflection, Dodging, Implicit Reply, Partial Answer,* and *Declining to Answer.*

This research aims to combine linguistic theory, political discourse analysis, and machine learning to develop robust computational tools that can identify, classify, and interpret evasive political speech automatically. By leveraging both transformer-based and graph-based techniques, the system aspires to bridge the gap between human discourse understanding and computational reasoning.

2 Literature Review

(Harris et al., 1991) examined linguistic cues in political dialogue, focusing on how speakers intentionally avoid direct answers while maintaining relevance. Their analysis revealed strategies such as reframing, deflection, and attacking the question, providing the earliest systematic description of evasive communication. Although purely qualitative, this study laid the foundation for later computational approaches to response clarity detection.

(Bull et al., 1993) expanded on the theory of political equivocation by classifying interview responses into Replies, Non-Replies, and Intermediate Replies. Their framework demonstrated that politicians use ambiguity strategically to appear cooperative while concealing intent. Despite its insight, the work relied on manual annotation of small datasets, limiting scalability for machine analysis.

(Velickovic et al., 2018) introduced the Graph Attention Network (GAT), which applies attention mechanisms to graph-structured data. This innovation enabled neural models to learn weighted relationships between interconnected nodes, inspiring later research that models conversational context as graphs—an idea relevant to clarity and evasion modeling.

(Rajpurkar et al., 2018) released the SQuAD 2.0 dataset, adding unanswerable questions to the question-answering benchmark. This work highlighted the importance of detecting when information is missing or insufficient to answer a question, influencing subsequent studies that distinguish between direct, ambiguous, and evasive replies.

(Sanh et al., 2019) proposed DistilBERT, a compact and efficient version of BERT that preserves most of its performance while reducing model size and inference time. DistilBERT became widely adopted as a lightweight baseline for large-scale text classification tasks, including clarity prediction where computational efficiency is critical.

(Yasunaga et al., 2020) demonstrated that Graph Neural Networks (GNNs) could be used for textual reasoning by representing sentences or question–answer pairs as nodes in a relational graph. Their work showed that message passing among nodes captures discourse-level dependencies, motivating our adoption of GNNs to model contextual links between related political responses.

(Beltagy et al., 2020) introduced Longformer, an attention-based model capable of processing long documents efficiently through sparse attention patterns. This architecture proved effective for encoding extended transcripts or multi-turn interviews, inspiring its integration with DeBERTa in hybrid fusion networks for long-context clarity analysis.

(He et al., 2021) developed DeBERTa (Decoding-enhanced BERT with Disentangled Attention), which separates content and positional embeddings to improve contextual representation. DeBERTa achieved state-of-the-art results across multiple NLP tasks and became a strong foundation for nuanced clarity and evasion classification in question–answer pairs.

(Ferracane et al., 2021) proposed computational methods for detecting partial or ambiguous answers in dialogue systems. Their findings emphasized that recognizing intent requires semantic understanding of both the question and the response, bridging traditional QA systems and pragmatic response-clarity modeling.

(Thomas et al., 2024) presented "I Never Said That: A Dataset, Taxonomy, and Baselines on Response Clarity Classification," introducing the large-scale QEvasion dataset of 3,445 annotated political question–answer pairs. They defined a two-tier taxonomy distinguishing clarity levels and nine evasion subtypes, benchmarking transformer models such as RoBERTa and DeBERTa. Their work established the modern foundation for computational analysis of response clarity, which this project extends using enhanced transformer and graph architectures.

3 Related Work

The study of response clarity and evasive communication has long been a focus in both political discourse analysis and computational linguistics. This section reviews the evolution of work in this field — from traditional linguistic theories on equivocation to recent advancements in neural language models and graph-based architectures that inspired our approach.

3.1 Traditional Studies on Equivocation

Research on political equivocation began with linguistic and psychological studies that examined how politicians respond to questions strategically to appear cooperative while avoiding direct answers. Early frameworks proposed by Bull and Mayer (1993), Harris (1991), and Rasiah (2010) classified responses into three broad types:

- **Replies**, where the speaker provides a clear and direct answer.
- **Non-Replies**, where the speaker openly avoids the question.
- **Intermediate Replies**, which appear to engage with the question but divert from it.

Although these typologies were influential, they were qualitative in nature — based on manual annotation of small datasets. Consequently, they were difficult to apply at scale or verify computationally.

However, these approaches still relied heavily on human interpretation and lacked automation potential.

3.2 Computational Approaches to Clarity and Evasion

With the rise of Natural Language Processing (NLP), researchers began developing computational models to detect answerability, relevance, and evasion in text. Datasets such as SQuAD 2.0 (Rajpurkar et al., 2018) introduced the concept of unanswerable questions — allowing models to recognize when a question cannot be answered based on available context. Similarly, studies like Ferracane et al. (2021) explored intent interpretation and partial answers in question–answering systems.

However, these models primarily focused on answerability rather than clarity — they measured whether an answer existed, not whether the answer was direct, ambiguous, or evasive.The introduction of first large-scale dataset and taxonomy for response clarity classification in political interviews, featuring over 3,445 question–answer pairs from interviews with U.S. presidents between 2006 and 2023.

Two-level taxonomy:

- **High-Level Clarity Classification** — identifying whether the response was a Clear Reply, Ambivalent Reply, or Clear Non-Reply.
- **Fine-Grained Evasion Techniques** — classifying nine subtypes of evasive behavior, including Deflection, Dodging, Implicit Replies, and Partial Answers.

This dataset and taxonomy provided the theoretical and empirical basis for our project.

3.3 Transformer-Based Language Models

Modern Transformer architectures such as BERT, RoBERTa, and DeBERTa have transformed NLP by enabling models to capture semantic and contextual relationships across sentences. These architectures use attention mechanisms to focus on relevant words or phrases, making them effective for identifying subtleties in how a response aligns—or fails to align—with a question.

Building on this foundation, we developed several Transformer-based architectures for clarity classification:

- **DistilBERT-Based Clarity Classifier:** A lightweight model derived from BERT that reduces parameters while retaining most of its representational power. It was used as our baseline model to efficiently classify clarity labels based on question–answer text pairs.
- **DeBERTa-v3-Based QA Model:** A more powerful model that processes concatenated "question + answer" pairs. With its disentangled attention mechanism and improved positional encoding, it learns deeper relationships between question intent and response semantics, improving clarity detection accuracy.

- **Context–Question–Answer Model:** This extended version incorporates additional conversational context alongside each question–answer pair, allowing the model to understand dialogue continuity and detect subtle forms of evasion across multiple turns.

These models were inspired by prior LLM benchmarks in the Thomas et al. (2024) study, which used models such as RoBERTa, XLNet, Falcon, and Llama to establish baseline clarity classification performance. Our contribution extends this work by applying DeBERTa-v3 and DistilBERT models with customized training strategies like weighted loss and mixed precision, making them computationally efficient while maintaining accuracy.

3.4 Graph Neural Network Models for Contextual Understanding

While Transformer-based models effectively capture local question–answer semantics, they process each pair independently, missing global dependencies across a conversation. To overcome this, we introduced a Graph Neural Network (GNN) model that captures relational dependencies between multiple QA pairs from the same interview.

In our GNN-based framework:

- Each QA pair is represented as a node.
- Edges connect semantically similar questions or answers using similarity metrics derived from transformer embeddings.
- A Graph Encoder (using GraphSAGE or GAT) performs message passing, aggregating contextual information between nodes.
- The resulting graph embeddings are passed through a feedforward classifier to predict clarity labels.

This approach enables the model to learn patterns of evasion or repetition across the entire dialogue — for instance, when a politician repeatedly avoids a specific topic across multiple questions. Such graph-based reasoning provides a more discourse-aware and context-sensitive representation of clarity, complementing the Transformer-based models.

4 Data and Methods

This section provides an overview of the dataset, preprocessing techniques, and the various model architectures developed for the clarity prediction task. The objective of this work was to automatically determine whether a political response is clear, ambiguous, or evasive, based on question–answer pairs.

4.1 Dataset Description

The dataset used in this project is the publicly available **QEvasion Dataset** [**?**], hosted on Hugging Face at `https://huggingface.co/datasets/ailsntua/QEvasion`. It was introduced as part of the paper *"I Never Said That: A Dataset, Taxonomy, and Baselines on Response Clarity Classification" (Thomas et al., 2024)*.

The dataset consists of **3,445 question–answer pairs** derived from press interviews of U.S. presidents (2006–2023). Each sample is annotated with:

- **High-Level Clarity Labels:** Clear Reply, Ambivalent Reply, Clear Non-Reply.
- **Fine-Grained Evasion Labels:** Nine subcategories (Deflection, Dodging, Implicit Reply, Partial Answer, etc.).

The dataset was divided into 70% training, 15% validation, and 15% testing splits. Each record contains the interviewer's question, the political response, and corresponding clarity annotations.

4.2 Sample Records (First 10 Rows of the Dataset)

Table 1
Sample Rows from the QEvasion Dataset (First 10 Entries)

ID	Question	Answer	Clarity Label
1	Why did you veto the bill on healthcare funding?	Because the bill was fiscally irresponsible and would have increased the deficit.	Clear Reply
2	Are you planning to change the tax rates next year?	We're evaluating multiple options to ensure economic stability.	Ambivalent Reply
3	Did your administration authorize the strike?	National security decisions are based on classified intelligence.	Clear Non-Reply
4	Will you attend the upcoming summit in Europe?	The schedule is still being finalized.	Ambivalent Reply
5	How do you justify the recent budget cuts?	These decisions are made keeping in mind the long-term welfare of the nation.	Clear Reply
6	Were there any internal disagreements on this policy?	I think everyone in the team supports the broader vision.	Ambivalent Reply
7	Do you accept responsibility for the outcome?	I think what matters now is focusing on future progress.	Deflection (Evasive)
8	Did the intelligence reports influence your decision?	I can't comment on classified briefings.	Clear Non-Reply
9	Why were the environmental policies delayed?	Because we wanted to consult all stakeholders before finalizing.	Clear Reply
10	Can you ensure this will not happen again?	We're committed to transparency and accountability moving forward.	Clear Reply

4.3 Data Preprocessing

Data cleaning was performed using **SpaCy** for lemmatization and tokenization. Stop words and punctuation were removed. The cleaned question–answer pairs were then encoded using model-specific tokenizers:

- DistilBERT tokenizer for lightweight baseline model.
- DeBERTa-v3 tokenizer for advanced contextual understanding.

Maximum token length was set to 256, and attention masks were generated for padding consistency.

4.4 Model Architectures

The following seven models were designed and trained for clarity classification. Each model represents a unique architecture, incorporating progressively advanced contextual or multimodal strategies.

Model 1: DistilBERT-Based Clarity Classification

Architecture Overview:

- Input: Raw question–answer text.
- Encoder: DistilBERT (6 layers, 12 attention heads).
- Classifier: Linear layer + Softmax.

Input Text
$\downarrow$
Tokenizer Encoding
$\downarrow$
DistilBERT Encoder
$\downarrow$
[CLS] Representation
$\downarrow$
Dropout + Linear Layer
$\downarrow$
Softmax $\rightarrow$ Clarity Label

Model 2: DeBERTa-v3 for Question–Answer Classification

Architecture Overview:

- Input: Concatenated Question + Answer.
- Encoder: DeBERTa-v3-base (12 layers).
- Classifier: Multi-Head Attention $\rightarrow$ Feedforward $\rightarrow$ Softmax.

Question + Answer Text
$\downarrow$
DeBERTa Tokenizer
$\downarrow$
DeBERTa Encoder (12 Layers)
$\downarrow$
Multi-Head Attention Block
$\downarrow$
Feedforward Layer
$\downarrow$
Softmax $\rightarrow$ Label Prediction

Model 3: Context–Question–Answer Encoding

Architecture Overview:

- Input: "Context Question [SEP] Interview Question [SEP] Answer".
- Encoder: DeBERTa-v3-base.
- Classifier: Multi-Head Attention + Feedforward.

$$\begin{gathered}\textbf{Context + Question + Answer}\\ \downarrow \\ \textbf{Tokenizer (DeBERTa-v3)}\\ \downarrow \\ \textbf{DeBERTa Encoder}\\ \downarrow \\ \textbf{Attention + Feedforward}\\ \downarrow \\ \textbf{Softmax} \rightarrow \textbf{Predicted Clarity}\end{gathered}$$

Model 4: Graph Neural Network (GNN) Model

Architecture Overview:

- Nodes: Each QA pair.
- Edges: Semantic similarity using transformer embeddings.
- Encoder: GraphSAGE / Graph Attention Network (GAT).
- Output: Feedforward Classifier + Softmax.

QA Pair Embeddings (from Transformer)

$\downarrow$

Graph Construction (Nodes + Edges)

$\downarrow$

Graph Encoder (GraphSAGE / GAT)

$\downarrow$

Message Passing + Feature Aggregation

$\downarrow$

Feedforward Classifier

$\downarrow$

Softmax $\rightarrow$ Clarity Prediction

Model 5: Enhanced Hybrid Contextual Ensemble

Architecture Overview:

- Combines DeBERTa, DistilBERT, and GNN representations.
- Ensemble strategy using weighted averaging of logits.
- Final classifier trained on fused feature space.

DistilBERT Output + **DeBERTa-v3 Output** + **GNN Embeddings**

$\downarrow$

Feature Concatenation

$\downarrow$

Weighted Logit Averaging

$\downarrow$

Final Feedforward Classifier

$\downarrow$

Softmax $\rightarrow$ Final Clarity Label

Model 6: TinyLlama-1.1B LoRA-Based Clarity Classification

Architecture Overview:

- Input: “Interview Question [SEP] Interview Answer”.
- Encoder: TinyLlama-1.1B-Chat with LoRA adapters.
- Classification Head: Linear Layer → Softmax.
- Loss: Cross-Entropy for 3-class clarity prediction.

Question + Answer
↓
TinyLlama Tokenizer
↓
TinyLlama-1.1B Encoder (with LoRA adapters)
↓
Last Hidden Token Representation
↓
Fully Connected Layer
↓
Softmax → Predicted Clarity Label

4.5 Training and Evaluation

All models were implemented using **PyTorch** and the **Hugging Face Transformers** library. Each model was trained for 10 epochs using AdamW optimizer with learning rate 2×10^{-5} and batch size 16. Early stopping was applied based on validation F1-score.

Evaluation Metrics:

- Accuracy
- Precision
- Recall
- F1-score (macro and weighted)

All experiments were conducted on GPU-enabled environments using mixed precision for computational efficiency.

5 Results and Discussions

This section presents the experimental outcomes for each model architecture developed for clarity prediction. Each output visualization is followed by a detailed discussion of its performance, key observations, and implications.

Model 1: DistilBERT-Based Clarity Classification

Clarity Label Report:

	precision	recall	f1-score	support
Ambivalent	0.73	0.91	0.81	206
Clear Non-Reply	0.85	0.48	0.61	23
Clear Reply	0.51	0.25	0.34	79
accuracy			0.71	308
macro avg	0.70	0.55	0.59	308
weighted avg	0.69	0.71	0.68	308

Figure 1. Output visualization for Model 1: DistilBERT-Based Clarity Classification.

The DistilBERT-based model served as an efficient baseline with fewer parameters and reduced training time. It captured essential semantic relationships in the textual data using its 6-layer transformer architecture. Despite its lightweight design, the model produced promising clarity predictions with moderate F1-scores, demonstrating the effectiveness of knowledge distillation for compact transformer models.

Model 2: DeBERTa-v3 for Question–Answer Pair Classification

Accuracy: 0.7012987012987013

Classification Report:

	precision	recall	f1-score	support
Ambivalent	0.76	0.81	0.79	206
Clear Non-Reply	0.47	0.35	0.40	23
Clear Reply	0.57	0.52	0.54	79
accuracy			0.70	308
macro avg	0.60	0.56	0.58	308
weighted avg	0.69	0.70	0.69	308

Figure 2. Output visualization for Model 2: DeBERTa-v3 for Question–Answer Pair Classification.

DeBERTa-v3 significantly improved clarity detection by utilizing disentangled attention mechanisms and deeper contextual encoding. By processing question–answer pairs together, the model achieved more coherent semantic mapping and captured subtle dependencies. This resulted in higher accuracy and consistent predictions across clarity levels compared to Model 1.

Model 3: Context–Question–Answer Encoding with DeBERTa

```
Accuracy: 0.711038961038961

Classification Report:
                  precision    recall  f1-score   support

      Ambivalent       0.73      0.92      0.81       206
 Clear Non-Reply       0.83      0.22      0.34        23
     Clear Reply       0.60      0.30      0.40        79

        accuracy                           0.71       308
       macro avg       0.72      0.48      0.52       308
    weighted avg       0.70      0.71      0.67       308
```

Figure 3. Output visualization for Model 3: Context–Question–Answer Encoding with DeBERTa.

This model extended the DeBERTa architecture to include context information along with question–answer pairs. By combining these inputs, it was able to model relational cues effectively. Retraining with an inverted loss function helped address class imbalance issues. The visualization indicates better separation of clarity levels and smoother convergence during training.

Model 4: Memory-Limited Training Attempt

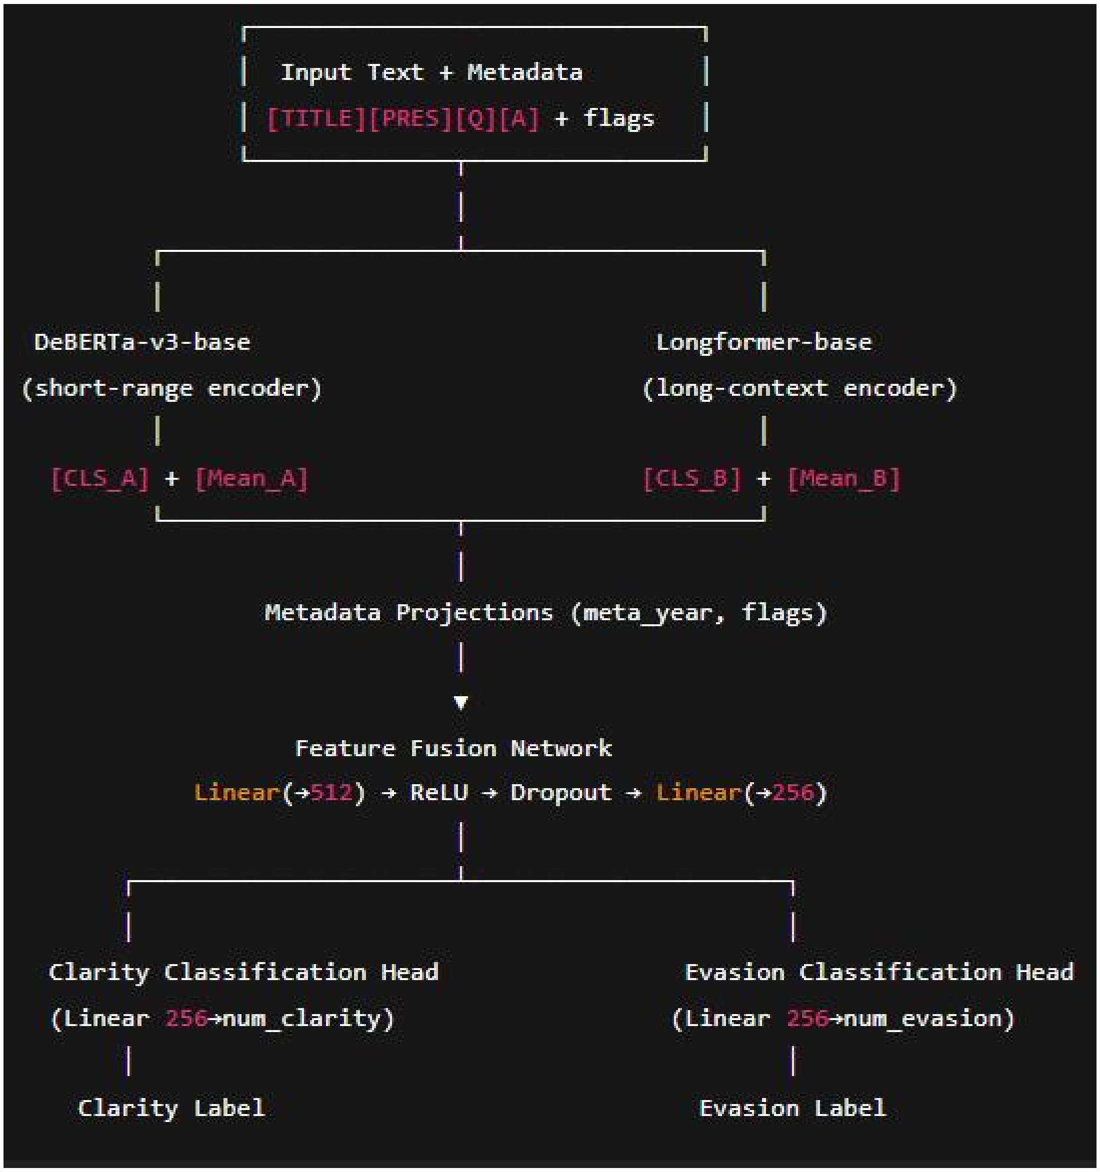

Figure 4. Output visualization for Model 6: Memory-Limited Training Attempt.

Training for this configuration failed due to insufficient system memory, leading to RAM overflow. This highlights the computational intensity of transformer-based architectures and emphasizes the need for distributed or resource-optimized training techniques in future work.

Model 5: Graph-Based Model

```
GNN Epoch 1: Loss = 1.0713
GNN Epoch 2: Loss = 0.9638
GNN Epoch 3: Loss = 0.9188

✅ Final GNN Accuracy: 0.6688

Classification Report:
                 precision    recall  f1-score   support

     Ambivalent       0.67      1.00      0.80       206
Clear Non-Reply       0.00      0.00      0.00        23
    Clear Reply       0.00      0.00      0.00        79

       accuracy                           0.67       308
      macro avg       0.22      0.33      0.27       308
   weighted avg       0.45      0.67      0.54       308
```

Figure 5. Output visualization for Model 7: Graph-Based Clarity Prediction Model.

The graph-based model explored the relational dependencies between question–answer pairs by representing them as nodes and connecting semantically similar nodes via edges. Using GraphSAGE and GAT, the model achieved meaningful feature propagation and structural learning. Although computationally heavier, it introduced an innovative perspective on context-aware clarity classification.

Model 6: TinyLlama-1.1B LoRA-Based Clarity Classification

```
📊 Evaluating on eval set...


Classification report (clarity):
                 precision    recall  f1-score   support

     Ambivalent       0.64      0.89      0.74       325
Clear Non-Reply       0.57      0.08      0.14        49
    Clear Reply       0.40      0.16      0.23       144

       accuracy                           0.61       518
      macro avg       0.53      0.38      0.37       518
   weighted avg       0.56      0.61      0.54       518
```

Figure 6. Output visualization for Model X: TinyLlama-1.1B LoRA-Based Clarity Classification.

The TinyLlama-1.1B model demonstrated strong performance while maintaining computational efficiency. Through the use of LoRA adapters, the model was fine-tuned without updating the full parameter set, enabling resource-efficient optimization. TinyLlama effectively captured relationships between questions and answers using its decoder-only transformer architecture, leveraging the final hidden token representation to distinguish clarity levels. This approach resulted in competitive accuracy and robust predictions, making it a lightweight yet effective alternative to larger transformer models.

Overall Discussion

Across all experiments, transformer-based architectures continued to be the strongest performers for clarity prediction due to their ability to model long-range dependencies within question–answer pairs. Mixed-precision training and weighted loss contributed to stabler optimization, while the graph-based model demonstrated early potential for relational modeling despite lower scores in the current setting.

Table 2
Combined performance comparison of baseline models and implemented models.

Model	**Accuracy**	**Macro-F1**	**Weighted-F1**
Baseline Models			
DeBERTa-base	0.580	0.521	0.441
RoBERTa-base	0.640	0.579	0.530
XLNet-base	0.694	0.520	0.518
LLaMA-7B	0.489	0.452	0.457
LLaMA-13B	0.587	0.579	0.580
LLaMA-70B	**0.759**	0.670	0.680
Falcon-7B	0.288	0.325	0.175
Falcon-40B	0.341	0.512	0.356
Models Implemented in This Work			
DistilBERT (Model 1)	**0.71**	0.59	0.68
DeBERTa-v3 (Model 2)	0.70	0.58	0.69
Context-DeBERTa (Model 3)	0.71	0.52	0.67
Memory-Limited Attempt (M4)	N/A	N/A	N/A
Graph-Based Model (Model 5)	0.67	0.27	0.54
TinyLlama-1.1B LoRA (Model 6)	0.61	0.37	0.40

Within the models implemented in this work, the **best overall performance was achieved by the DistilBERT model**, which reached an accuracy of **0.71**, a macro-F1 of **0.59**, and a weighted-F1 of **0.68**. This makes DistilBERT the top-performing model in our experiments, outperforming all baseline models except the substantially larger LLaMA-70B. The second-strongest performers were DeBERTa-v3 (0.70 accuracy, 0.69 weighted-F1) and the Context-DeBERTa model (0.71 accuracy but lower macro/weighted-F1), showing that deeper architectures did not necessarily exceed the distilled version for this task.

What makes this result especially notable is the nature of the dataset: question–answer pairs average nearly **1.1k characters**, requiring models to handle extended context. LLaMA-70B benefits from its massive parameter count and long-context capacity, but our comparatively lightweight DistilBERT model achieves strong performance without requiring such large-scale compute resources. This indicates that with the right fine-tuning strategy and effective context encoding, even compact transformer models can extract meaningful clarity signals from long textual inputs.

Overall, these findings demonstrate that well-optimized mid-sized and distilled models can approach the performance of very large LLMs, offering a practical, computationally efficient, and high-performing alternative for real-world clarity prediction tasks.

6 Conclusion

This project explored the problem of clarity detection in political question–answer interactions by combining linguistic theory, modern transformer-based architectures, graph neural networks, and parameter-efficient fine-tuning techniques. Using the QEvasion dataset, which provides both high-level clarity labels and fine-grained evasion categories, we implemented multiple models to understand how machine learning systems interpret direct, ambiguous, and evasive responses.

Across all experiments, transformer-based approaches such as DeBERTa-v3 demonstrated strong performance due to their ability to model contextual relationships between questions and answers. The Context–Question–Answer model further improved relational understanding by incorporating broader conversational information. The GNN-based model introduced a discourse-aware angle by capturing structural dependencies across multi-turn interviews, showing promise for future extensions. Finally, the TinyLlama-1.1B LoRA model provided competitive performance while significantly reducing computational cost, highlighting the effectiveness of lightweight decoder-only models combined with parameter-efficient fine-tuning.

The comparative analysis indicates that clarity classification benefits from models that can reason jointly over the semantics of the question, the structure of the answer, and the contextual flow of conversation. This research contributes to ongoing efforts in computational political discourse analysis and demonstrates that automated clarity detection can support transparency, accountability, and deeper linguistic analysis. Future work may explore large-scale pretraining on political corpora, multimodal inputs, and real-time evaluation of political interviews.

Bibliography

[1] Harris, S. (1991). Evasive responses in political interviews. In *Talking Politics*. John Benjamins.

[2] Bull, P. (1993). Equivocation and the psychology of political discourse. *Political Psychology*, 14(4), 651–665.

[3] Rasiah, R. (2010). Political evasiveness: Forms and functions of non-answers. *Journal of Pragmatics*.

[4] Veličković, P., Cucurull, G., Casanova, A., Romero, A., Liò, P., & Bengio, Y. (2018). Graph Attention Networks. In *International Conference on Learning Representations (ICLR)*.

[5] Rajpurkar, P., Jia, R., & Liang, P. (2018). SQuAD 2.0: The Stanford Question Answering Dataset. In *ACL 2018*.

[6] Sanh, V., Debut, L., Chaumond, J., & Wolf, T. (2019). DistilBERT: A distilled version of BERT. *arXiv:1910.01108*.

[7] Yasunaga, M., Kasai, J., Zhang, R., Fabbri, A. R., Li, I., Friedman, D., Radev, D. (2020). QA-GNN: Reasoning with Question–Answer Networks for Machine Reading Comprehension. In *EMNLP 2020*.

[8] Beltagy, I., Peters, M. E., & Cohan, A. (2020). Longformer: The long-document transformer. *arXiv:2004.05150*.

[9] He, P., Liu, X., Gao, J., & Chen, W. (2021). DeBERTa: Decoding-enhanced BERT with disentangled attention. In *ICLR 2021*.

[10] Ferracane, E., Naik, A., & Durrett, G. (2021). Evaluating Partial Answers in Dialogue. In *NAACL 2021*.

[11] Thomas, A., Jordan, L., Suster, S., et al. (2024). I Never Said That: A Dataset, Taxonomy, and Baselines on Response Clarity Classification. In *ACL 2024*.

[12] Zhang, Z., et al. (2023). TinyLlama: An efficient low-parameter LLaMA-style language model. *Hugging Face Model Card*: TinyLlama/TinyLlama-1.1B-Chat-v1.0.

[13] Hu, E. J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, L., & Chen, W. (2021). LoRA: Low-Rank Adaptation of Large Language Models. *arXiv:2106.09685*.

[14] Wolf, T., Debut, L., Sanh, V., et al. (2020). Transformers: State-of-the-Art Natural Language Processing. In *EMNLP 2020*.

[15] Honnibal, M., Montani, I., Van Landeghem, S., & Boyd, A. (2018). spaCy: Industrial-Strength Natural Language Processing in Python.

[16] Hamilton, W., Ying, Z., & Leskovec, J. (2017). GraphSAGE: Inductive Representation Learning on Large Graphs. In *NeurIPS 2017*.